

# ML\_Assignment\_5

Sushmitha\_Virri

#700742289

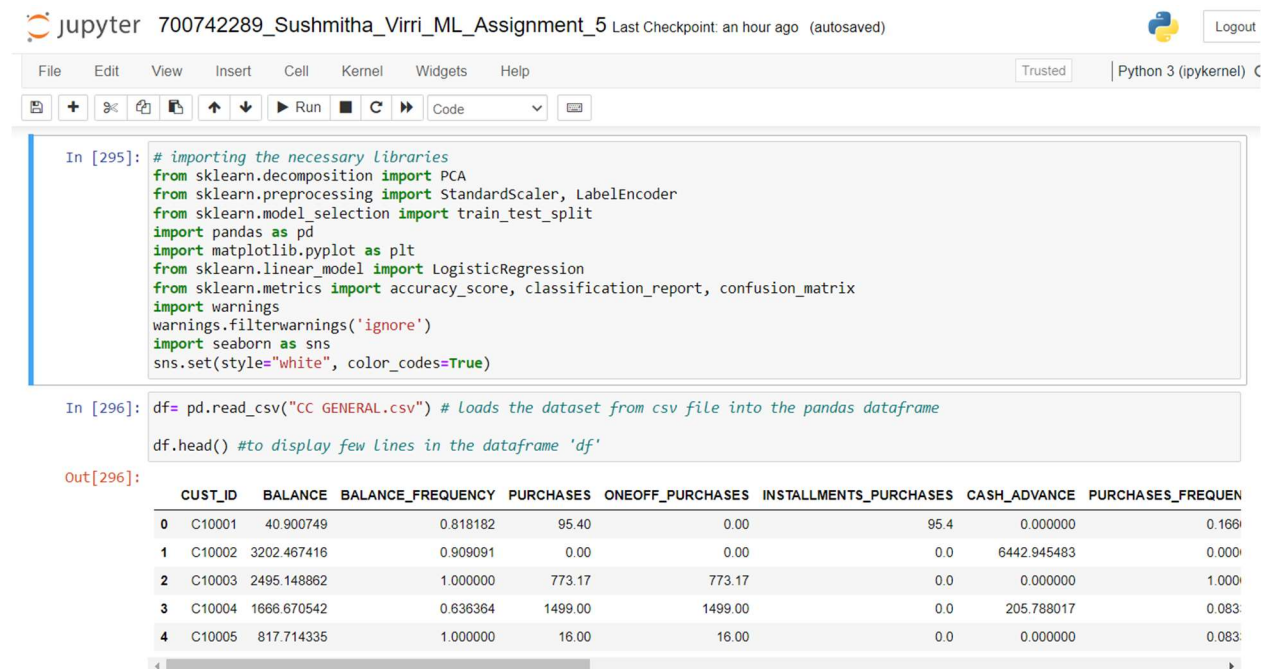
Github Link: <https://github.com/Sushmitha->

[Virri/MLAssignments21627/blob/main/700742289\\_Sushmitha\\_Virri\\_ML\\_Assignment\\_5.ipynb](https://github.com/Sushmitha-Virri/MLAssignments21627/blob/main/700742289_Sushmitha_Virri_ML_Assignment_5.ipynb)

Drive Video Link:

<https://drive.google.com/file/d/10XSx5hnwqFJPNDWGBte3bJkWGIMZHB0X/view?usp=sharing>

## 1. Principal Component Analysis



The screenshot shows a Jupyter Notebook titled "700742289\_Sushmitha\_Virri\_ML\_Assignment\_5". The interface includes a top bar with the Jupyter logo, the notebook title, and a "Logout" button. Below the top bar is a menu bar with options: File, Edit, View, Insert, Cell, Kernel, Widgets, Help. A toolbar with icons for file operations and a "Run" button is also present. The notebook content area shows two code cells. The first cell, labeled "In [295]:", contains a block of Python code that imports various libraries: sklearn.decomposition (PCA), sklearn.preprocessing (StandardScaler, LabelEncoder), sklearn.model\_selection (train\_test\_split), pandas (pd), matplotlib.pyplot (plt), sklearn.linear\_model (LogisticRegression), sklearn.metrics (accuracy\_score, classification\_report, confusion\_matrix), warnings, and seaborn (sns). It also sets warnings to be ignored and configures the seaborn style to "white" with color codes. The second cell, labeled "In [296]:", contains code to load a CSV file named "CC\_GENERAL.csv" into a pandas DataFrame named 'df' and to display the first few rows of the DataFrame. The output of the second cell, labeled "Out[296]:", shows a preview of the DataFrame with columns: CUST\_ID, BALANCE, BALANCE\_FREQUENCY, PURCHASES, ONEOFF\_PURCHASES, INSTALLMENTS\_PURCHASES, CASH\_ADVANCE, PURCHASES\_FREQUEN, and an unlabeled column with numerical values.

```
In [295]: # importing the necessary libraries
from sklearn.decomposition import PCA
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.model_selection import train_test_split
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
import warnings
warnings.filterwarnings('ignore')
import seaborn as sns
sns.set(style="white", color_codes=True)

In [296]: df = pd.read_csv("CC_GENERAL.csv") # Loads the dataset from csv file into the pandas dataframe
df.head() #to display few lines in the dataframe 'df'
```

Out[296]:

	CUST_ID	BALANCE	BALANCE_FREQUENCY	PURCHASES	ONEOFF_PURCHASES	INSTALLMENTS_PURCHASES	CASH_ADVANCE	PURCHASES_FREQUEN
0	C10001	40.900749	0.818182	95.40	0.00	95.4	0.000000	0.166
1	C10002	3202.467416	0.909091	0.00	0.00	0.0	6442.945483	0.000
2	C10003	2495.148862	1.000000	773.17	773.17	0.0	0.000000	1.000
3	C10004	1666.670542	0.636364	1499.00	1499.00	0.0	205.788017	0.083
4	C10005	817.714335	1.000000	16.00	16.00	0.0	0.000000	0.083

In the above python script, we have imported the required modules from the libraries, including pandas for data manipulation, StandardScaler for feature scaling, PCA for dimensionality reduction, KMeans for clustering and metrics evaluation.

We have loaded the CC GENEREAL dataset from csv file to pandas data frame.

head() function displays the first few lines of the dataframe.

Next we use df.isnull().any() function to check whether there are any missing values. If any, fill them with mean of each column using fillna() function.

‘inplace = True’ argument will make changes to the original dataframe and is displayed in the output

File Edit View Insert Cell Kernel Widgets Help








 Run
 


 Code
 


In [297]: `df.shape # using shape attribute to get the shape of df`

Out[297]: (8950, 18)

In [298]: `df.isnull().any()`

Out[298]:

CUST_ID	False
BALANCE	False
BALANCE_FREQUENCY	False
PURCHASES	False
ONEOFF_PURCHASES	False
INSTALLMENTS_PURCHASES	False
CASH_ADVANCE	False
PURCHASES_FREQUENCY	False
ONEOFF_PURCHASES_FREQUENCY	False
PURCHASES_INSTALLMENTS_FREQUENCY	False
CASH_ADVANCE_FREQUENCY	False
CASH_ADVANCE_TRX	False
PURCHASES_TRX	False
CREDIT_LIMIT	True
PAYMENTS	False
MINIMUM_PAYMENTS	True
PRC_FULL_PAYMENT	False
TENURE	False
dtype:	bool

In [299]: `df.fillna(df.mean(), inplace = True)`  
`df.isnull().any()`

Out[299]:

CUST_ID	False
BALANCE	False
BALANCE_FREQUENCY	False
PURCHASES	False
ONEOFF_PURCHASES	False
INSTALLMENTS_PURCHASES	False
CASH_ADVANCE	False
PURCHASES_FREQUENCY	False

```
File Edit View Insert Cell Kernel Widgets Help
In [300]: x = df.iloc[:,[1,2,3,4]] # extracts 1,2,3,4 columns from the dataset and assign it to 'x'
          y = df.iloc[:, -1] # extracts the last column and assign it to 'y'
          print(x.shape, y.shape)
          (8950, 4) (8950,)
```

The above code selects a subset of columns at index 1,2,3,4 from pandas Dataframe 'df' and assigns it to 'x' and selects all rows and last column and assign it to 'y'.

Then it prints the shape of 'x' and 'y' which represent the number of rows and columns in the output cell.

### a. Apply PCA on CC dataset.

```
File Edit View Insert Cell Kernel Widgets Help Trusted
In [301]: # a. Apply PCA on CC GENERAL dataset.
          # PCA is used to analyse the high dimensional dataset
          # so as to reduce the complexity of the data and extract important features.

          pca = PCA(3) #creating a PCA with 3 principal components
          x_pca = pca.fit_transform(x) # fits the PCA object to the dataset 'x' then transforming it to new dataset 'x_pca'

          # creating a new pandas dataframe 'df2' with less number of dimensions.
          df2 = pd.DataFrame(data = x_pca, columns = ['principal component 1',
                                                    'principal component 2', 'principal component 3'])

          # Concatinates principalDF with last column of 'df' and create new dataframe 'df'
          df3 = pd.concat([df2, df.iloc[:, -1]], axis = 1)
          df3.head()

Out[301]:
```

	principal component 1	principal component 2	principal component 3	TENURE
0	-1500.250819	-1114.178407	-64.9889145	12
1	-592.910661	1914.657567	-151.542222	12
2	217.734556	905.144354	-291.615901	12
3	927.782551	-198.923616	-421.617772	12
4	-1310.548986	-359.591021	-132.892069	12

In this code we applied PCA on CC GENERAL dataset. In the first line we created a 'pca' object with 3 principal components.

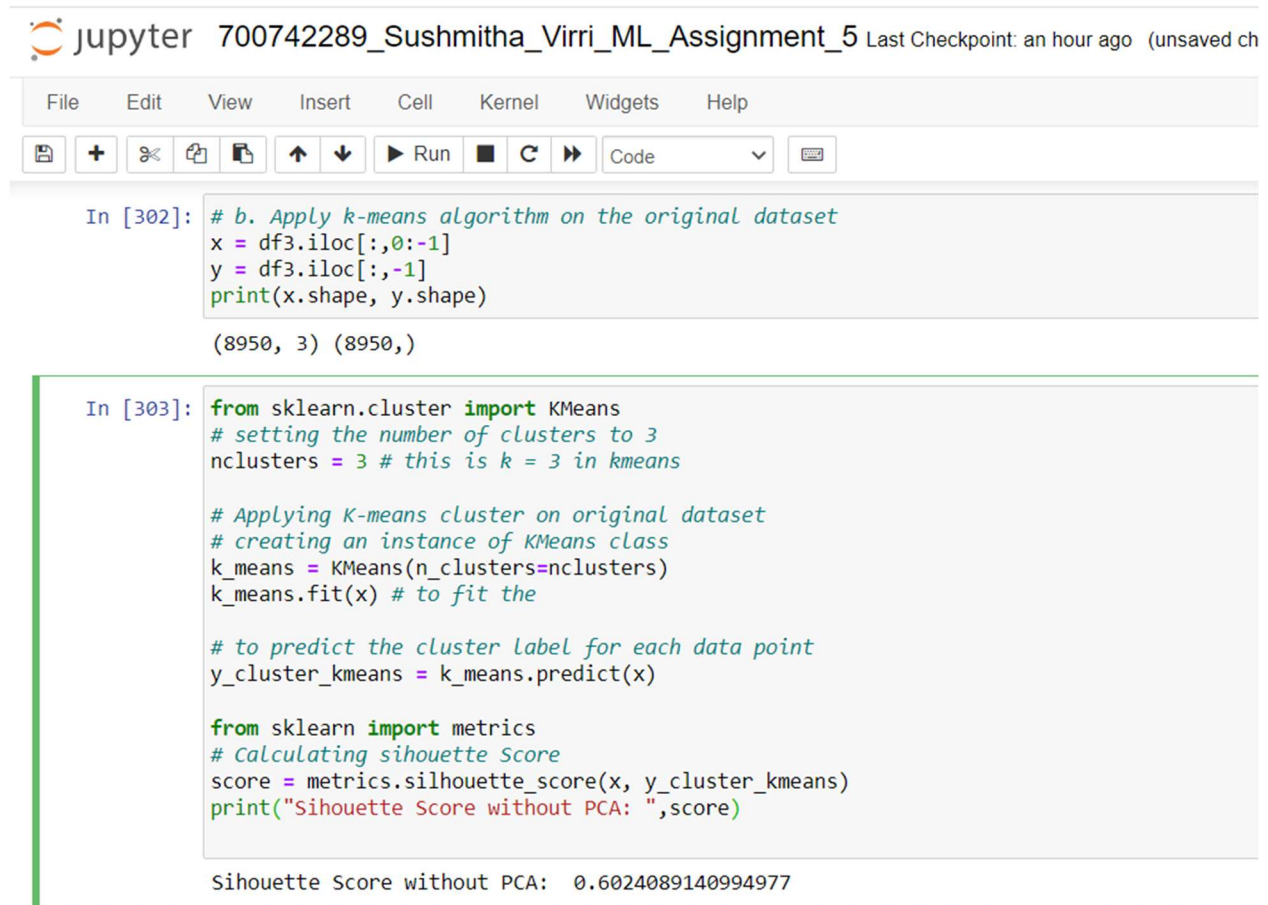
In the second line PCA object is fit to the original dataset 'x' and transformed to a new dataset 'x\_pca'

The third statement creates a new pandas dataframe 'df2' with principal component columns.

Next statement concatenates 'df2' with last column of 'df' and assigned to a new dataframe 'df3'

The output cell shows reduced dimensions of the dataset.

**b. Apply k-means algorithm on the PCA result and report your observation if the silhouette score has improved or not?**



```
In [302]: # b. Apply k-means algorithm on the original dataset
x = df3.iloc[:,0:-1]
y = df3.iloc[:,1]
print(x.shape, y.shape)

(8950, 3) (8950,)
```

```
In [303]: from sklearn.cluster import KMeans
# setting the number of clusters to 3
nclusters = 3 # this is k = 3 in kmeans

# Applying K-means cluster on original dataset
# creating an instance of KMeans class
k_means = KMeans(n_clusters=nclusters)
k_means.fit(x) # to fit the

# to predict the cluster label for each data point
y_cluster_kmeans = k_means.predict(x)

from sklearn import metrics
# Calculating silhouette Score
score = metrics.silhouette_score(x, y_cluster_kmeans)
print("Silhouette Score without PCA: ",score)

Silhouette Score without PCA: 0.6024089140994977
```

The above screenshot talks about the application of KMeans clustering algorithm from scikit-learn library to the original dataset.

The number of clusters is set to 3 and stored in variable nclusters. KMeans class instance is created with this value and fit() method is called to fit the data.

Predict() method is used to predict the cluster label for each datapoint.

silhouette\_score() function calculates the score with input values as original dataset and predicted labels.

```
File Edit View Insert Cell Kernel Widgets Help
[Icons] [Run] [Code]

In [304]: # Applying KMeans cluster on reduced dataset
kmeans_pca = KMeans(n_clusters = nclusters)
kmeans_pca.fit(x_pca)
y_cluster_kmeans_pca = kmeans_pca.predict(x_pca)

# Calculating silhouette Score
score_pca = metrics.silhouette_score(x_pca, y_cluster_kmeans_pca)
print("Silhouette Score with PCA: ",score_pca)

Silhouette Score with PCA: 0.6024089140994978
```

The above code applies KMeans clustering to preprocessed dataset with PCA.

An instance of KMeans class is created with nclusters. fit() method is called on this instance with input as 'x\_pca' and predict() method is then called on kmeans\_pca to predict the cluster labels for PCA transformed dataset x\_pca.

Silhouette\_score is then calculated with PCA-transformed dataset and cluster labels as input.

**Performance:** From the above we can observe that the silhouette score with PCA is slightly less than the silhouette score without PCA.

**Reason:** Implementing PCA has reduced the number of features in the dataset.

### c. Perform Scaling+PCA+K-Means and report performance.

To perform dimensionality reduction using PCA and clustering using K-means on the dataset.

The code for this creates an instance of StandardScaler class and fits it on the training data 'x' and then transforms it into fitted scaler.

The code then performs feature scaling using StandardScaler, dimensionality reduction using PCA, clustering using KMeans and then evaluates clustering performance using silhouette score.

```
File Edit View Insert Cell Kernel Widgets Help
In [305]: # StandardScaler is used to scale the features to have mean '0' and variance '1'
scaler = StandardScaler()
scaler.fit(x) # for fitting on training data 'x'
x_scale = scaler.transform(x)

pca2 = PCA(3)
x_pca2 = pca.fit_transform(x_scale)

df4 = pd.DataFrame(data = x_pca2, columns = ['principal component 1',
                                             'principal component 2', 'principal component 3'])

final_df = pd.concat([df4, df[['TENURE']]], axis = 1)
final_df.head()

from sklearn.cluster import KMeans
nclusters = 3 # this is the k in kmeans
km = KMeans(n_clusters=nclusters)
km.fit(x_scale)

# predict the cluster for each data point
y_cluster_kmeans = km.predict(x_scale)
from sklearn import metrics
score = metrics.silhouette_score(x_scale, y_cluster_kmeans)
print(score)
```

0.5573894918696056

### Performance report:

We can observe from the output that the silhouette score after applying Scaling + PCA + K-means is reduced when compared to the silhouette scores without PCA and with PCA.

From this we can say that as the silhouette score is not improved that is it is not greater than the previous values and hence the performance is not improved.



## 2. Use pd\_speech\_features.csv

### a. Perform Scaling

### b. Apply PCA (k=3)

### c. Use SVM to report performance

jupyter 700742289\_Sushmitha\_Virri\_ML\_Assignment\_5 Last Checkpoint: 3 hours ago (unsaved changes)

File Edit View Insert Cell Kernel Widgets Help

Run

```
In [306]: # to Load the dataset from csv file into the pandas dataframe
df_pd = pd.read_csv(r"pd_speech_features.csv")
```

```
In [307]: df_pd.head()
```

Out[307]:

	id	gender	PPE	DFA	RPDE	numPulses	numPeriodsPulses	meanPeriodPulses	stdDevPeriodPulses	locPctJitter	...
0	0	1	0.85247	0.71826	0.57227	240	239	0.008064	0.000087	0.00218	...
1	0	1	0.76686	0.69481	0.53966	234	233	0.008258	0.000073	0.00195	...
2	0	1	0.85083	0.67604	0.58982	232	231	0.008340	0.000060	0.00176	...
3	1	0	0.41121	0.79672	0.59257	178	177	0.010858	0.000183	0.00419	...
4	1	0	0.32790	0.79782	0.53028	236	235	0.008162	0.002669	0.00535	...

5 rows × 755 columns

```
In [308]: df_pd.isnull().any()
```

Out[308]:

id	False
gender	False
PPE	False
DFA	False
RPDE	False
...	...
tqwt_kurtosisValue_dec_33	False
tqwt_kurtosisValue_dec_34	False
tqwt_kurtosisValue_dec_35	False
tqwt_kurtosisValue_dec_36	False
class	False

Length: 755, dtype: bool



File Edit View Insert Cell Kernel Widgets Help

```
In [309]: x = df_pd.drop('class',axis=1).values
          y = df_pd['class'].values
          print(x.shape,y.shape)

          (756, 754) (756,)
```

```
In [310]: #a. to perform Scaling
          scaler = StandardScaler()
          x_scale = scaler.fit_transform(x)
          #b. Applying PCA for k=3
          pca = PCA(3)
          x_pca = pca.fit_transform(x_scale)

          principalDf = pd.DataFrame(data = x_pca, columns = ['principal component 1', 'principal component 2'
          finalDf = pd.concat([principalDf, df_pd[['class']]], axis = 1)
          finalDf.head()
```

Out[310]:

	principal component 1	principal component 2	Principal Component 3	class
0	-10.047372	1.471076	-6.846404	1
1	-10.637725	1.583749	-6.830976	1
2	-13.516185	-1.253543	-6.818700	1
3	-9.155084	8.833600	15.290905	1
4	-6.764470	4.611467	15.637124	1

a. Scaling:

In the first step we scale the data using StandardScaler() function and stored in the variable 'x\_scale'.

b. PCA:

The next step is to perform PCA on the scaled data using PCA() function with k=3 and store it in the variable 'x\_pca'.

Next we create a dataframe 'principalDf' with three principal components. The code then concatenates 'principalDf' with 'class' column using the pd.concat() function.



```
File Edit View Insert Cell Kernel Widgets Help
```

```

In [311]: #c. to use SVM and report performace

#splitting the data into training and testing data
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.3,random_state=0)

from sklearn.svm import SVC
svm_classifier = SVC()
svm_classifier.fit(x_train, y_train)

#predict() method of the class SVC is used to predict the target variable for the test set.
y_pred = svm_classifier.predict(x_test)

# Evaluations made by the classifier
print(classification_report(y_test, y_pred, zero_division=1))
print(confusion_matrix(y_test, y_pred))

# Accuracy score
acc_svc = accuracy_score(y_pred,y_test)
print('accuracy is',acc_svc)

#Calculate sihouette Score
score = metrics.silhouette_score(x_test, y_pred)
print("Sihouette Score: ",score)

```

	precision	recall	f1-score	support
0	1.00	0.02	0.03	57
1	0.75	1.00	0.86	170
accuracy			0.75	227
macro avg	0.88	0.51	0.45	227
weighted avg	0.81	0.75	0.65	227

```

[[ 1 56]
 [ 0 170]]
accuracy is 0.7533039647577092
Sihouette Score: 0.8052538192732682


```

### c. Support Vector Machine(SVM):

The data is split into training and testing sets using `train_test_split()` function. The SVM classifier is trained on the training set using the `fit()` method.

`classification_report()` and `confusion_matrix()` are used to determine the performance of the classifier.

`silhouette_score()` measures the similarity of data points within the cluster and the dissimilarity of data points among different clusters.



3. Apply Linear Discriminant Analysis (LDA) on Iris.csv dataset to reduce dimensionality of data to k=2

```
In [312]: import math
import numpy as np
df_iris = pd.read_csv(r"Iris.csv")
df_iris.head()
```

```
Out[312]:
```

	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	Species
0	1	5.1	3.5	1.4	0.2	Iris-setosa
1	2	4.9	3.0	1.4	0.2	Iris-setosa
2	3	4.7	3.2	1.3	0.2	Iris-setosa
3	4	4.6	3.1	1.5	0.2	Iris-setosa
4	5	5.0	3.6	1.4	0.2	Iris-setosa

```
In [313]: df_iris.isnull().any()
```

```
Out[313]: Id                False
SepalLengthCm             False
SepalWidthCm              False
PetalLengthCm             False
PetalWidthCm             False
Species                   False
dtype: bool
```

```
In [317]: x = df_iris.iloc[:,1:-1]
y = df_iris.iloc[:, -1]
print(x.shape,y.shape)
```

```
(150, 4) (150,)
```

The above code reads the data from 'iris.csv' file into the pandas dataframe 'df\_iris'. Displays the first five columns. Next checks for any missing values in the dataframe, then separates dataframe into 2 parts using 'iloc' method.

1. 'x' which contain all features of dataset except target variable.
2. 'y' which contain only the target variable.

```
File Edit View Insert Cell Kernel Widgets Help
[Icons] [Run] [Code]

In [320]: x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.3, random_state=0)
          from sklearn.preprocessing import StandardScaler

          le = LabelEncoder()
          y = le.fit_transform(y)
          # StandardScaler is used to standardize the data prior to modelling
          scaler = StandardScaler()
          x_train = scaler.fit_transform(x_train)
          x_test = scaler.transform(x_test)
          le = LabelEncoder()
          y = le.fit_transform(y)
          #fit_transform method calculates mean and standard deviation of every feature
          x_train_std = scaler.fit_transform(df.iloc[:,1:-1].values)

In [321]: from sklearn.preprocessing import LabelEncoder

          #LabelEncoder class encodes the categorical valued to numerical values
          le = LabelEncoder()
          y_le = le.fit_transform(df.iloc[:, -1].values)

In [328]: from sklearn.discriminant_analysis import LinearDiscriminantAnalysis as LDA
          lda = LDA(n_components=2)
          x_train = lda.fit_transform(x_train, y_train)
          x_test = lda.transform(x_test)
          print(x_train.shape, x_test.shape)

(105, 2) (45, 2)
```

The above code performs following operations:

Splitting the data into training and testing sets using train\_test\_split' function from the sklearn library.

Encoding the target variable using 'LabelEncoder'

Standardizing the target variable using 'StandardScaler'

Then applying (LDA) Linear Discriminant Analysis to reduce feature dimension to 2.

Thus the code does preprocessing and reduced the dimensions using LDA.

#### 4.Explain briefly the distinction between PCA and LDA.

LDA and PCA both use linear transformations to maximize variance in a smaller dimension. The PCA method is an unsupervised learning algorithm, whereas the LDA algorithm is a supervised learning system. This means that PCA seeks maximum variance directions regardless of class labels, whereas LDA finds maximum class separability directions.

LDA and PCA both use linear transformations to maximize variance in a lower dimension. The PCA method is an unsupervised learning algorithm, whereas the LDA algorithm is a supervised learning system. This indicates that PCA finds maximum variance directions regardless of class labels, whereas LDA finds maximum class separability directions.

PCA : It condenses the characteristics into a smaller set of orthogonal variables known as principal components, which are linear combinations of the original variables. The first component captures the most variability in the data, the second the second, and so on. It reduces the features to a smaller group of orthogonal variables called principal components - linear combinations of the original variables. The first component captures the most variability in the data, the second the second most, and so on.

LDA : LDA finds linear discriminants in order to maximize the variance between the different categories while minimizing the variance within the class.